

Intelligent Candidate Discovery & Ranking

Evidence-weighted scoring across 100,000 candidates for Senior AI Engineer — Founding Team

100K

Candidates Ranked

~60s

CPU Runtime

20

Honeypots Flagged

Mukul

mukul24359@iiitd.ac.in · IIIT Delhi, CSE 2028

github.com/Mukul07777/redrob-ranker

mukul001-redrob-ranker.hf.space

Why Keyword Filters Fail

✗ The Trap

A Marketing Manager lists 'FAISS', 'RAG', 'Pinecone' as skills and passes every keyword filter.

✗ The Blind Spot

A Tier-5 candidate who shipped a production recommender at Swiggy never mentions 'vector database' — and gets filtered out.

✗ The JD Warning

The JD explicitly states: 'The right answer is NOT finding candidates whose skills section contains the most AI keywords.'

Evidence Over Keywords

✓ Evidence-Weighted Skills

Skill score = proficiency × endorsements × duration. Keywords with zero evidence get 5% trust regardless of claimed level.

✓ Career Context

Title trajectory, product vs consulting company, ML signals in job descriptions — not just the skills list.

✓ Behavioral Availability

A perfect-on-paper candidate inactive for 7 months with 5% response rate is not actually hireable.

✓ Honeypot Immunity

4 mathematical impossibility checks eliminate fabricated profiles before scoring begins.

7-Component Weighted Formula

$\text{final_score} = \text{skill}(38\%) \times w + \text{career}(22\%) + w + \text{availability}(15\%) + w + \text{yoe}(12\%) + w + \text{location}(7\%) + w + \text{notice}(4\%) + w + \text{github}(2\%)$

● 38%

Skill Match

proficiency $\times$ endorsement trust $\times$ duration trust. Core JD skills (embeddings, vector DB, ranking, eval) get 3 $\times$ weight.

● 22%

Career Quality

Product company bonus.
TCS/Infosys/Wipro/Accenture/Cognizant/Capgemini = heavy penalty. ML signals in job descriptions add score.

● 15%

Availability

Recency of last login $\times$ open_to_work $\times$ recruiter response rate $\times$ interview completion. Inactive >6 months $\rightarrow$ max 0.35 multiplier.

● 12%

Experience Years

Sweet spot 6–8 yrs = 1.0.
Symmetrical ramp-down. >12 yrs penalized for title-chaser pattern per JD.

● 7%

Location

Noida/Pune/Delhi NCR/Hyderabad = 1.0. Tier-1 + willing to relocate = 0.85. Outside India (no visa) = 0.55.

● 4%

Notice Period

<30 days = 1.0. 30–60 = 0.85. 60–90 = 0.72. 90+ = 0.60. JD explicitly prefers sub-30 day notice periods.

● 2%

GitHub Activity

Open-source proxy. No GitHub = 0.4 neutral-negative. Scales linearly to 1.0 at activity score 80+.

Honeypot Detection — 4 Impossibility Checks

⚠ ~80 honeypots exist in the 100K pool · Submissions with >10% honeypots in top-100 are DISQUALIFIED · All 20 flagged candidates scored 0.001

01

YoE vs Career History Mismatch

FLAGGED → score = 0.001

Claimed years_of_experience is $>3\times$ the implied YoE from summed career_history duration_months. Example: claims 15 years but all job entries total 4 years.

02

Expert Skills With Zero Evidence

FLAGGED → score = 0.001

4+ skills marked 'advanced' or 'expert' with BOTH endorsements=0 AND duration_months=0. Legitimate experts always leave at least one evidence trail.

03

Single Job + Impossible Tenure

FLAGGED → score = 0.001

Claims 10+ years experience, has exactly one job entry, and started that job <6 months ago. Mathematically impossible — would require pre-birth employment.

04

Endorsement-Duration Ratio Anomaly

FLAGGED → score = 0.001

3+ skills each with >50 endorsements but ≤ 2 months duration. Endorsements accumulate over years — high endorsements on brand-new skills signal fabrication.

Top 5 Ranked Candidates

✓ **Submission Valid**

python validate_submission.py → Submission is valid.

20 Honeypots

Detected and suppressed from top 100

~60 Seconds

Full 100K ranking on CPU, stdlib only

1

Staff Machine Learning Engineer

CAND_0077337 · Swiggy · Delhi NCR · 7.0 yrs

Embeddings, vector search, ranking · 91% response rate · Active · <30d · Notice <30d **0.95**
39

2

Lead AI Engineer

CAND_0081846 · Product Co · Preferred City · 6.7 yrs

Core retrieval + NLP skills with endorsements · High availability · Strong behavioral signals **0.93**
63

3

Senior Machine Learning Engineer

CAND_0018499 · Product Co · India · 7.2 yrs

Production ML background · JD sweet-spot YoE · Strong platform engagement **0.93**
27

4

Senior AI Engineer

CAND_0002025 · AI-Native Co · India · 5.9 yrs

Search/ranking experience · Endorsed skills with duration · Good location fit **0.93**
22


5

Search Engineer

CAND_0064326 · Product Co · India · 7.6 yrs

Direct role title match · IR + ranking background · Active and responsive **0.92**
65

How the Skill Match Score Works



WHY THIS BEATS KEYWORD MATCHING → A candidate claiming 'Expert' in FAISS with 0 endorsements and 0 months scores lower than one with 'Intermediate' + 40 endorsements + 36 months.

Catching Keyword Stuffers

The Marketing Manager with perfect AI skills scores near 0 on career_quality regardless of keyword count.

Consulting-only career (TCS/Infosys entire history) → automatic 0.35 penalty on career score per JD instructions.

Elevating Tier-5 Candidates

Text bonus scans candidate summary + headline for core concepts — even if not listed as formal skills.

A Swiggy engineer who describes 'built a real-time recommendation engine' scores a text bonus for 'recommendation' even without 'FAISS' in skill list.

Everything Delivered



submission.csv

100 candidates ranked, validated by validate_submission.py



rank.py

Single-file ranker — stdlib only, CPU, no network, ~60s runtime



Sandbox

mukul001-redrob-ranker.hf.space — live Gradio demo



GitHub Repo

github.com/Mukul07777/redrob-ranker — public, clean, documented



PDF Deck

This presentation — architecture, approach, results



submission_metadata.yaml

Team info, AI tools declared, reproduce command provided